

```
In [1053]: import numpy as np
import pandas as pd
# For visualizations
import matplotlib.pyplot as plt
import seaborn as sns
from mycolorpy import colorlist as mcp

# For regular expressions
import re
# For handling string
import string
# For performing mathematical operations
import math

from sklearn.metrics import silhouette_score
from scipy.stats import zscore
```

## Exploratory Data Analysis

\*\*Authored by: Sangita Baitalik

```
In [1054]: #An automobile company has plans to enter new markets with their existing products.
#After intensive market research, they've deduced that the behavior of new markets is different from their existing market.
#In their existing market, the sales team has classified all customers into 4 segments.
#Then, they performed segmented outreach and communication for different segments.
#This strategy has work exceptionally well for them. They plan to use the same strategy for the new markets.
#You are required to help the manager to predict the right group of the new customers.
```

```
In [1055]: data=pd.read_csv("../files/customersegmentation.csv")
```

```
In [1056]: data.head()
```

```
Out[1056]:
```

	ID	Gender	Ever_Married	Age	Graduated	Profession	Work_Experience	Spending_Score
0	458982	Male	Yes	61	Yes	Executive	1.0	1.0
1	458983	Female	Yes	63	Yes	Executive	0.0	1.0
2	458984	Male	Yes	39	Yes	Artist	0.0	Average
3	458985	Male	No	23	No	Healthcare	1.0	1.0
4	458986	Male	No	18	No	Healthcare	7.0	1.0

```
In [1057]: data.shape
```

```
Out[1057]: (10695, 11)
```

```
In [1058]: # Looking for missing values in dataset
```

```
data.isna().sum()
```

```
Out[1058]: ID                0
           Gender            0
           Ever_Married      190
           Age               0
           Graduated         102
           Profession        162
           Work_Experience   1098
           Spending_Score    0
           Family_Size       448
           Var_1             108
           Segmentation      0
           dtype: int64
```

```
In [1059]: data = data.dropna()
           data.shape
```

```
Out[1059]: (8819, 11)
```

```
In [1060]: data.isna().sum()
```

```
Out[1060]: ID                0
           Gender            0
           Ever_Married      0
           Age               0
           Graduated         0
           Profession        0
           Work_Experience   0
           Spending_Score    0
           Family_Size       0
           Var_1             0
           Segmentation      0
           dtype: int64
```

### Gender Data Visualisation

```
In [1061]: data['Gender'].dtype
```

```
Out[1061]: dtype('O')
```

```
In [1062]: data['Gender'].unique()
```

```
Out[1062]: array(['Male', 'Female'], dtype=object)
```


```
In [1063]: data['Gender'].value_counts()
```

```
Out[1063]: Male      4861
           Female    3958
           Name: Gender, dtype: int64
```

```
In [1064]: labels=data['Gender'].unique()
           values=data['Gender'].value_counts(ascending=True)

fig, (ax0,ax1) = plt.subplots(ncols=2,figsize=(15,8))
bar = ax0.bar(x=labels, height=values, width=0.4, align='center', color=['pink','blue'])
ax0.set(title='Count difference in Gender Distribution',xlabel='Gender', ylabel='No. of Customers')
ax0.set_ylim(0,130)
ax0.axhline(y=data['Gender'].value_counts()[0], color='pink', linestyle='--', label='Female (4861)')
ax0.axhline(y=data['Gender'].value_counts()[1], color='blue', linestyle='--', label='Male (3958)')
ax0.legend()

ax1.pie(values,labels=labels,colors=['pink','blue'],autopct='%1.1f%%')
ax1.set(title='Ratio of Gender Distribution')
fig.suptitle('Gender Distribution', fontsize=20);
plt.show()
```



```
In [1065]: data.Gender=pd.Categorical(data.Gender, categories=['Male', 'Female'], ordered=True)
```

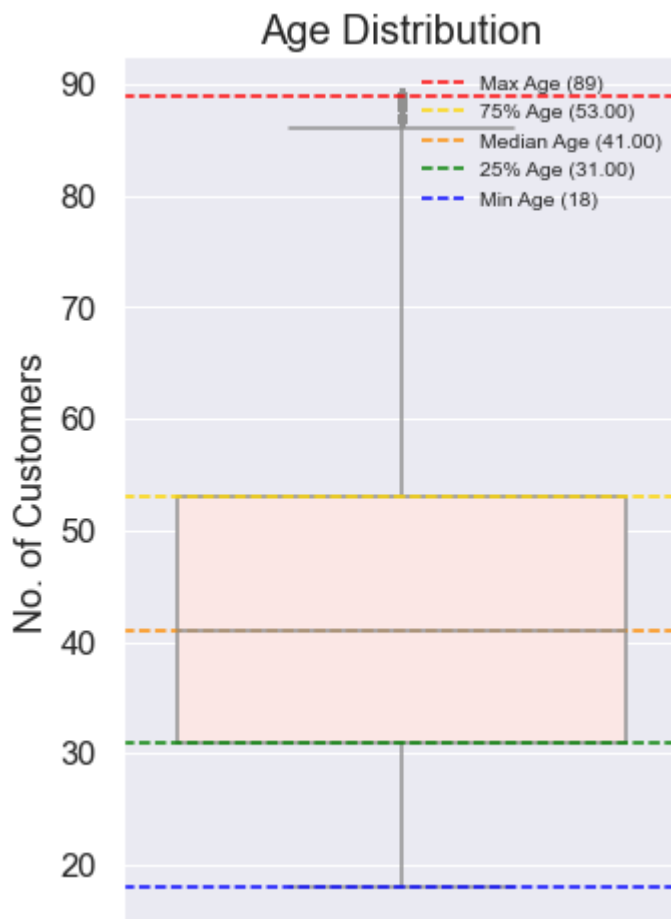
### Age Data Visualisation

```
In [1066]: data['Age'].describe()
```

```
Out[1066]: count      8819.000000
mean        43.517859
std         16.581537
min         18.000000
25%         31.000000
50%         41.000000
75%         53.000000
max         89.000000
Name: Age, dtype: float64
```


```
In [1067]: fig, ax = plt.subplots(figsize=(5,8))
sns.set(font_scale=1.5)
ax = sns.boxplot(y=data["Age"], color="mistyrose")
ax.axhline(y=data['Age'].max(), linestyle='--',color='red', label=f'Max Age (89)')
ax.axhline(y=data['Age'].describe()[6], linestyle='--',color='gold', label=f'75% Age (53.00)')
ax.axhline(y=data['Age'].median(), linestyle='--',color='darkorange', label=f'Median Age (41.00)')
ax.axhline(y=data['Age'].describe()[4], linestyle='--',color='green', label=f'25% Age (31.00)')
ax.axhline(y=data['Age'].min(), linestyle='--',color='blue', label=f'Min Age (18)')
ax.legend(fontsize='xx-small', loc='upper right')
ax.set_ylabel('No. of Customers')

plt.title('Age Distribution', fontsize = 20)
plt.show()
```



```
In [1068]: fig, ax = plt.subplots(figsize=(30,8))
sns.set(font_scale=1.5)
ax = sns.countplot(x=data['Age'], palette='cool')
ax.axhline(y=data['Age'].value_counts().max(), linestyle='--',color='darkblue')
ax.axhline(y=data['Age'].value_counts().mean(), linestyle='--',color='mediumblue')
ax.axhline(y=data['Age'].value_counts().min(), linestyle='--',color='slateblue')
ax.legend(loc='upper right')
ax.set_ylabel('No. of Customers')

plt.title('Age Distribution', fontsize = 20)
plt.show()
```

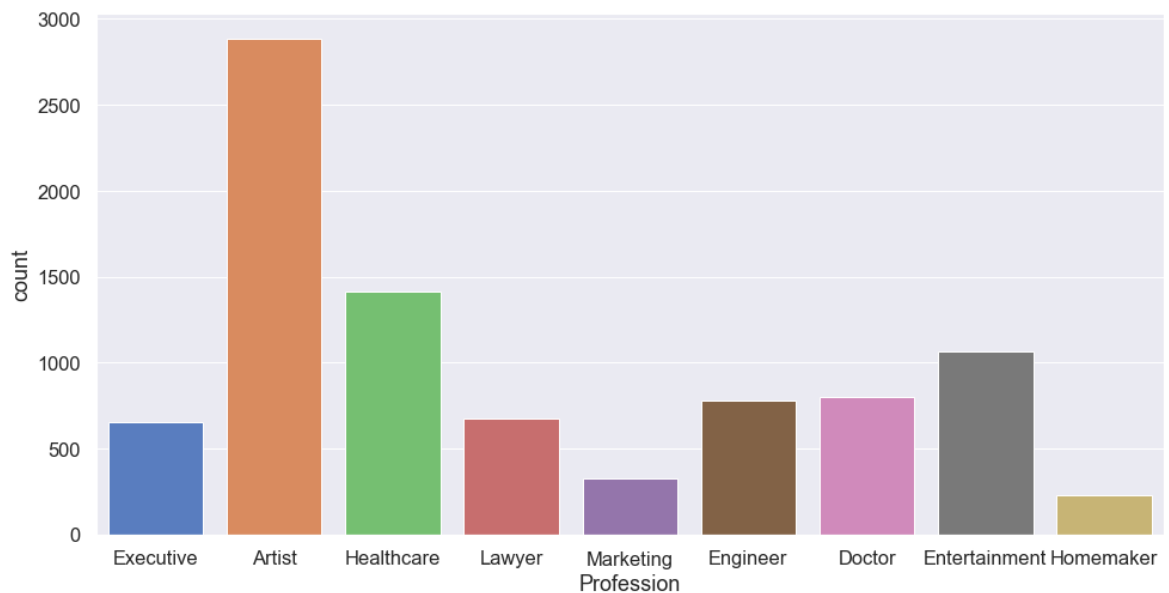


## Profession Data Visualisation

```
In [1069]: ▶ plt.figure(figsize=(16,8))  
sns.countplot(data.Profession,palette='muted')
```

C:\Users\shrut\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

Out[1069]: <AxesSubplot:xlabel='Profession', ylabel='count'>



```
In [1070]: ▶ profession=pd.get_dummies(data.Profession)  
data.drop(['Profession'],axis=1,inplace=True)  
data=data.join(profession)
```


Graduated Data Visualisation

```
In [1071]: sns.countplot(data.Graduated,palette='rocket')
```

C:\Users\shrut\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[1071]: <AxesSubplot:xlabel='Graduated', ylabel='count'>
```



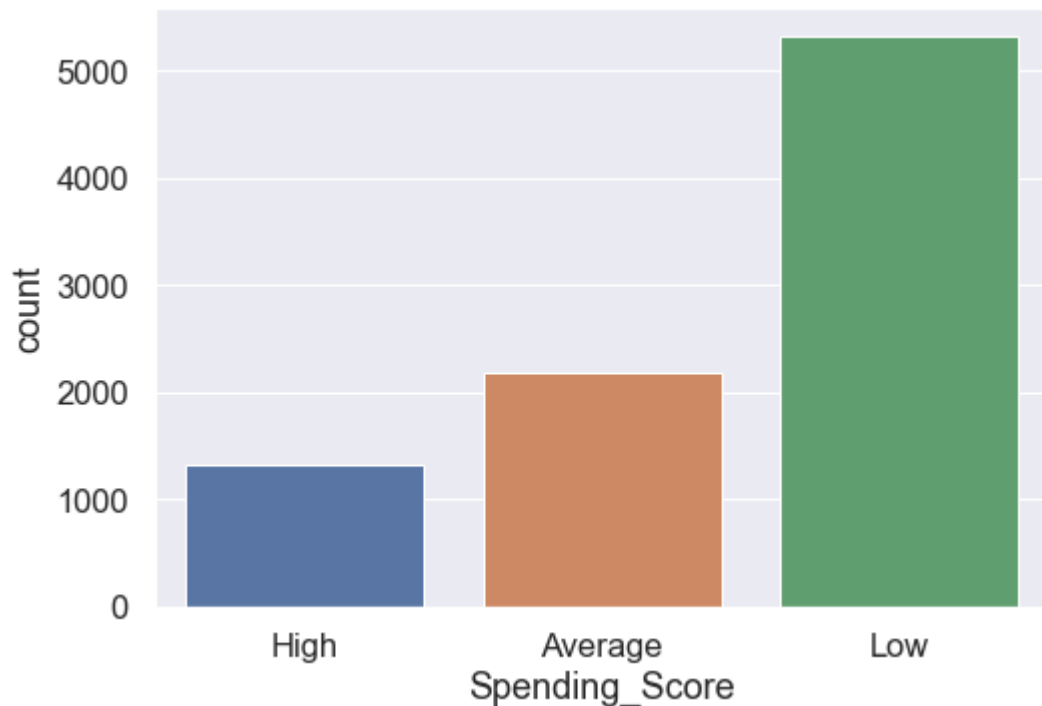
```
In [1072]: data.Graduated=pd.Categorical(data.Graduated,categories=['No','Yes'],ordered=
```

### Spending Score Visualisation

```
In [1073]: sns.countplot(data.Spending_Score)
```

```
C:\Users\shrut\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
```

```
Out[1073]: <AxesSubplot:xlabel='Spending_Score', ylabel='count'>
```



```
In [1074]: data.Spending_Score=pd.Categorical(data.Spending_Score,categories=['Low','Ave
```

Var\_1 Visualisation



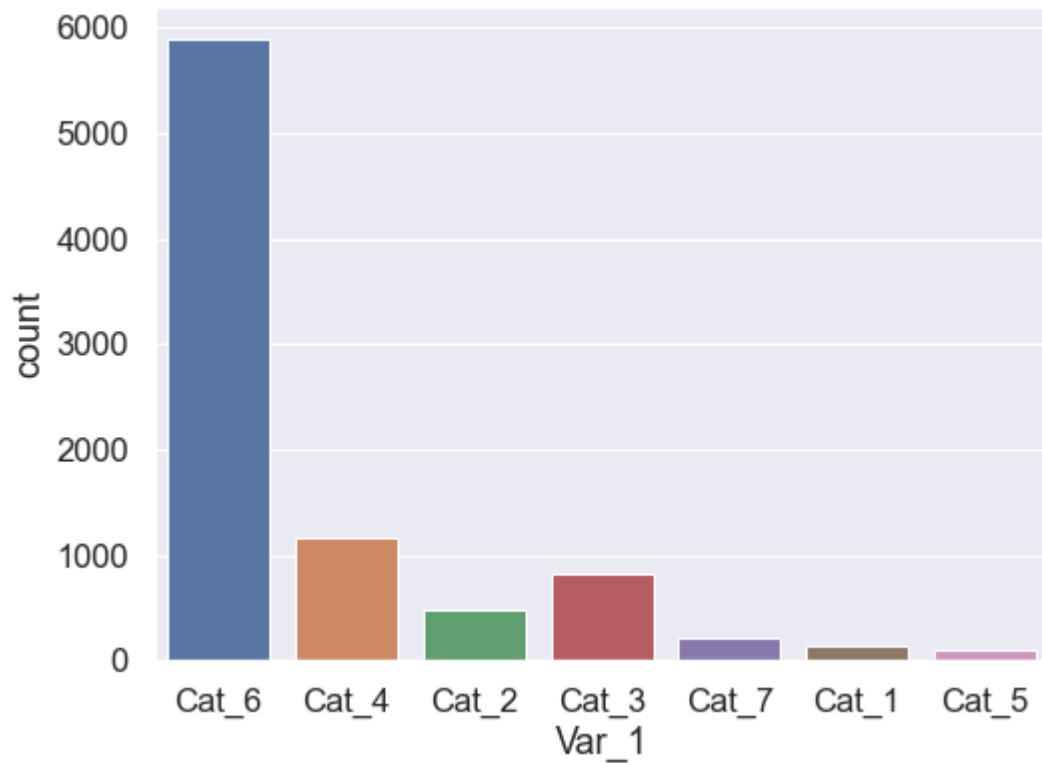
In [1075]: `#Var_1 is income range attribute with cat_1 being the highest paid and cat_6`

```
plt.figure(figsize=(8,6))  
sns.countplot(data.Var_1)
```

C:\Users\shrut\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[1075]: `<AxesSubplot:xlabel='Var_1', ylabel='count'>`



```
In [1076]: data.Var_1=pd.Categorical(data.Var_1).codes
```


### Marital Status Data Visualisation

```
In [1077]: sns.countplot(data.Ever_Married,palette='cubehelix')
```

C:\Users\shrut\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

```
Out[1077]: <AxesSubplot:xlabel='Ever_Married', ylabel='count'>
```




```
In [1078]: data.Ever_Married=pd.Categorical(data.Ever_Married,categories=['No','Yes'],or
```

### Work Experience Data Visualisation

```
In [1079]: plt.figure(figsize=(8,6))  
plt.hist(data.Work_Experience)
```

```
Out[1079]: (array([5729., 337., 631., 248., 242., 771., 569., 121., 55.,  
116.]),  
array([ 0. , 1.4, 2.8, 4.2, 5.6, 7. , 8.4, 9.8, 11.2, 12.6, 14. ]),  
<BarContainer object of 10 artists>)
```




```
label=pd.Categorical(data.Segmentation,categories=['A','B','C','D']).codes  
data.drop(['Segmentation'],axis=1,inplace=True) label
```

```
In [1080]: correlation_data=pd.DataFrame(data)
correlation_data.drop(['ID'],axis=1,inplace=True)
correlation_data
```

```
Out[1080]:
```


	Gender	Ever_Married	Age	Graduated	Work_Experience	Spending_Score	Family_Size
0	0	1	61	1	1.0	2	3.0
1	1	1	63	1	0.0	2	5.0
2	0	1	39	1	0.0	1	3.0
3	0	0	23	0	1.0	0	4.0
4	0	0	18	0	7.0	0	4.0
...	...	...	...	...	...	...	..
10690	1	1	43	1	0.0	1	2.0
10691	1	0	31	1	1.0	0	4.0
10692	0	0	22	0	1.0	0	3.0
10693	1	1	66	1	0.0	1	3.0
10694	1	0	43	1	1.0	0	1.0

8819 rows × 18 columns



```
In [1081]: plt.figure(figsize=(25,20))
sns.heatmap(correlation_data.corr(),annot=True)
```

Out[1081]: <AxesSubplot:>



## K-means Algorithm and Analysis

\*\* Authored by: Shruti Chanda

Calculate the z-score and remove the outliers from the dataset.

```
In [1082]: # Calculate z-score from the correlation except the
z_scores = zscore(correlation_data.drop('Segmentation', axis=1))


# Filtering rows with zscore less than 3
abs_z_scores = np.abs(z_scores)
filtered_entries = (abs_z_scores < 3).all(axis=1)
new_df = correlation_data[filtered_entries]
```

```
In [1083]: new_df
```

```
Out[1083]:
```

	Gender	Ever_Married	Age	Graduated	Work_Experience	Spending_Score	Family_Size
2	0	1	39	1	0.0	1	3.0
3	0	0	23	0	1.0	0	4.0
4	0	0	18	0	7.0	0	4.0
13	0	1	38	1	8.0	1	4.0
14	0	1	37	1	8.0	1	4.0
...	...	...	...	...	...	...	..
10689	1	0	43	1	9.0	0	3.0
10690	1	1	43	1	0.0	1	2.0
10691	1	0	31	1	1.0	0	4.0
10692	0	0	22	0	1.0	0	3.0
10694	1	0	43	1	1.0	0	1.0

5236 rows × 18 columns



### K-means Algorithm Class

```

In [1084]: # Generic class to fit and predict k-means clustering model
## k: the number of clusters desired
## tol: value helps identify model convergence
## max_iter: maximum number of iterations incase, model does not converges

class k_means:
    # Intitalize model parameters
    def __init__(self, k=2, tol=0.001, max_iter=300):
        self.k = k
        self.tol = tol
        self.max_iter = max_iter

    # Describes the relationship between a response variable and one or more
    def model_fit(self, data):
        # Create an empty liost for centroids
        self.centroids = {}

        # Randomly select k points to begin
        for i in range(self.k):
            self.centroids[i] = data[np.random.choice(range(len(data)), 1, re

        # Use initial centroids to label data rows to various k vales, then c
        for i in range(self.max_iter):
            self.classifications = {}

            for i in range(self.k):
                self.classifications[i] = []

            for features in data:
                distances = [math.sqrt(np.linalg.norm(features-self.centroids
                classification = distances.index(min(distances))
                self.classifications[classification].append(features)

            prev_centroids = dict(self.centroids)

            for classification in self.classifications:
                self.centroids[classification] = np.average(self.classificati

            opt = True

            for c in self.centroids:
                orig_centroid = prev_centroids[c]
                curr_centroid = self.centroids[c]
                if np.sum((curr_centroid-orig_centroid)/orig_centroid*100.0)
                    opt = False

            if opt:
                return opt

    # Used to predict outcomes by analyzing patterns in a given set of input
    def model_predict(self, data):
        classification = []

        # Returns the classification list
        for d_row in range(len(data)):
            distances = [math.sqrt(np.linalg.norm(data[d_row]-self.centroids[

```

```
classification.append(distances.index(min(distances)))
```

```
return classification
```

### Plotting functions

```
In [1085]: ▶ # Graph plotting function for k-means clustering which can be extended to k=1
def plot_graph (centroids_data, classification_data, k=2):
    color_lst = mcp.gen_color(cmap="cividis", n=k)
    mark_lst = ['x', '^', '>', '<', '8', 's', 'p', 'h', 'H', 'd', 'D']

    fig = plt.figure(figsize=(20, 20))
    ax = fig.add_subplot(111, projection='3d')

    # Plot cluster centroids
    for centroid in centroids_data:
        ax.scatter(centroids_data[centroid][0], centroids_data[centroid][1],
                  marker="o", color="blue", s=150, linewidths=5)

    # Plot classification data
    for classification in classification_data:
        color = color_lst[classification]
        mark = mark_lst[classification]
        for featureset in classification_data[classification]:
            ax.scatter(featureset[0], featureset[1], featureset[2], marker=mark, color=color)

    plt.show()

# Plotting cost functions obtained to analyze optimal value of k
def plot_costs(costs, trials=2):
    x = np.arange(2, trials)
    plt.plot(x, costs)
    plt.title("Elbow curve")
    plt.xlabel("K -->")
    plt.ylabel("Dispersion")

# Plotting the silhoutte plot to analyze best value of k
def plot_silhoutte(scores, trials=2):
    x = np.arange(2, trials)
    plt.plot(x, scores)
    plt.title("Silhoutte Score for k's")
    plt.xlabel("K -->")
    plt.ylabel("Score")
```

### Independent Helper Functions



```

In [1086]: # Generic function to compute Principal Component Analysis (PCA) for dimensi
def principalComponentAnalysis(data, n_components):

    # Mean centering the data
    X_mean = data - np.mean(data , axis = 0)

    # Calculating the covariance matrix of the mean-centered data.
    cov_mat = np.cov(X_mean , rowvar = False)

    # Calculating Eigenvalues and Eigenvectors of the covariance matrix
    eigen_vals , eigen_vecs = np.linalg.eigh(cov_mat)

    # Sort the eigenvalues in descending order
    sorted_vec = np.argsort(eigen_vals)[::-1]

    sorted_eigenvalue = eigen_vals[sorted_vec]

    # Similarly sort the eigenvectors
    sorted_eigenvectors = eigen_vecs[:,sorted_vec]

    # Select the first n eigenvectors, n is desired dimension of our final re
    eigenvector_subset = sorted_eigenvectors[:, 0:n_components]

    # Transform the data
    X_reduced = np.dot(eigenvector_subset.transpose(),X_mean.transpose()).tra

    return X_reduced

# Calculating cost function for various values of k
def cost_function(data, trials=1):
    costs = []
    scores = []
    # Run the loop for the number of trials
    for i in range(2,trials):
        # Initialize K means with different values of k
        kmeans = k_means(k=i)
        kmeans.model_fit(data)

        cluster_assignments = kmeans.centroids

        # Calculate the distance from their respective centroides for evaluat
        cost = 0
        for cluster in cluster_assignments:
            for feature in kmeans.classifications[cluster]:
                dist = np.linalg.norm(feature - cluster_assignments[cluster])
                cost += dist
            costs.append(np.array(cost))

        # Calculate the silhoutte score
        scores.append(silhouette_score(data, kmeans.model_predict(data), metr

    # Return cost and scores arrays
    return costs, scores

# Split a dataset into a train and test set
def train_test_split(df, frac=0.2):

```

```
# get random sample
test = df.sample(frac=frac, axis=0)

# get everything but the test sample
train = df.drop(index=test.index)

return train, test
```


### **Data Preparation**

```
In [1087]: ▶ # Split data into training and testing data
train, test = train_test_split(new_df, frac=0.6)
```

```
In [1088]: ▶ # Apply model fit on the training data
train_x = train.drop('Segmentation', axis=1)
train_x = train[['Age', 'Graduated', 'Work_Experience', 'Spending_Score']].va
train_y = train.Segmentation
train_y = pd.Categorical(train_y, categories=['A', 'B', 'C', 'D'], ordered=True).c
```


### **Analyze for best value of k**

```
In [1098]: # Implement PCA since we have 4 features  
updated_x = principalComponentAnalysis(train_x, 3)  
  
cost_arr, scores_arr = cost_function(updated_x, trials=10)  
  
# Generate cost plot for various values of k  
plot_costs(cost_arr, trials=10)
```



Using the Elbow curve we can see that the optimal value for k is 4 and can try to vary it between 4 to 6.


```
In [1099]: # Generate silhouette plot for various values of k  
plot_silhouette(scores_arr, trials=10)
```



Using the Silhouette score we can observe that we get the highest score with k=4.

```
In [1104]: ▶ # Apply model fit on the training data using the optimal value of k as 4
train_kmeans = k_means(k=4)
train_kmeans.model_fit(updated_x)

# Generate k-means graph for identified clusters
plot_graph(train_kmeans.centroids, train_kmeans.classifications, k=4)
```



```
In [1105]: ▶ # Predict cluster label outcomes  
out_y = train_kmeans.model_predict(updated_x)  
  
# Analyze the model performance  
score = silhouette_score(updated_x, out_y, metric='euclidean')  
  
print(f'Silhoutte Score: {score}%)
```

Silhoutte Score: 0.4531733863745631%

### ***Varying parameters to optimize model performance***

```

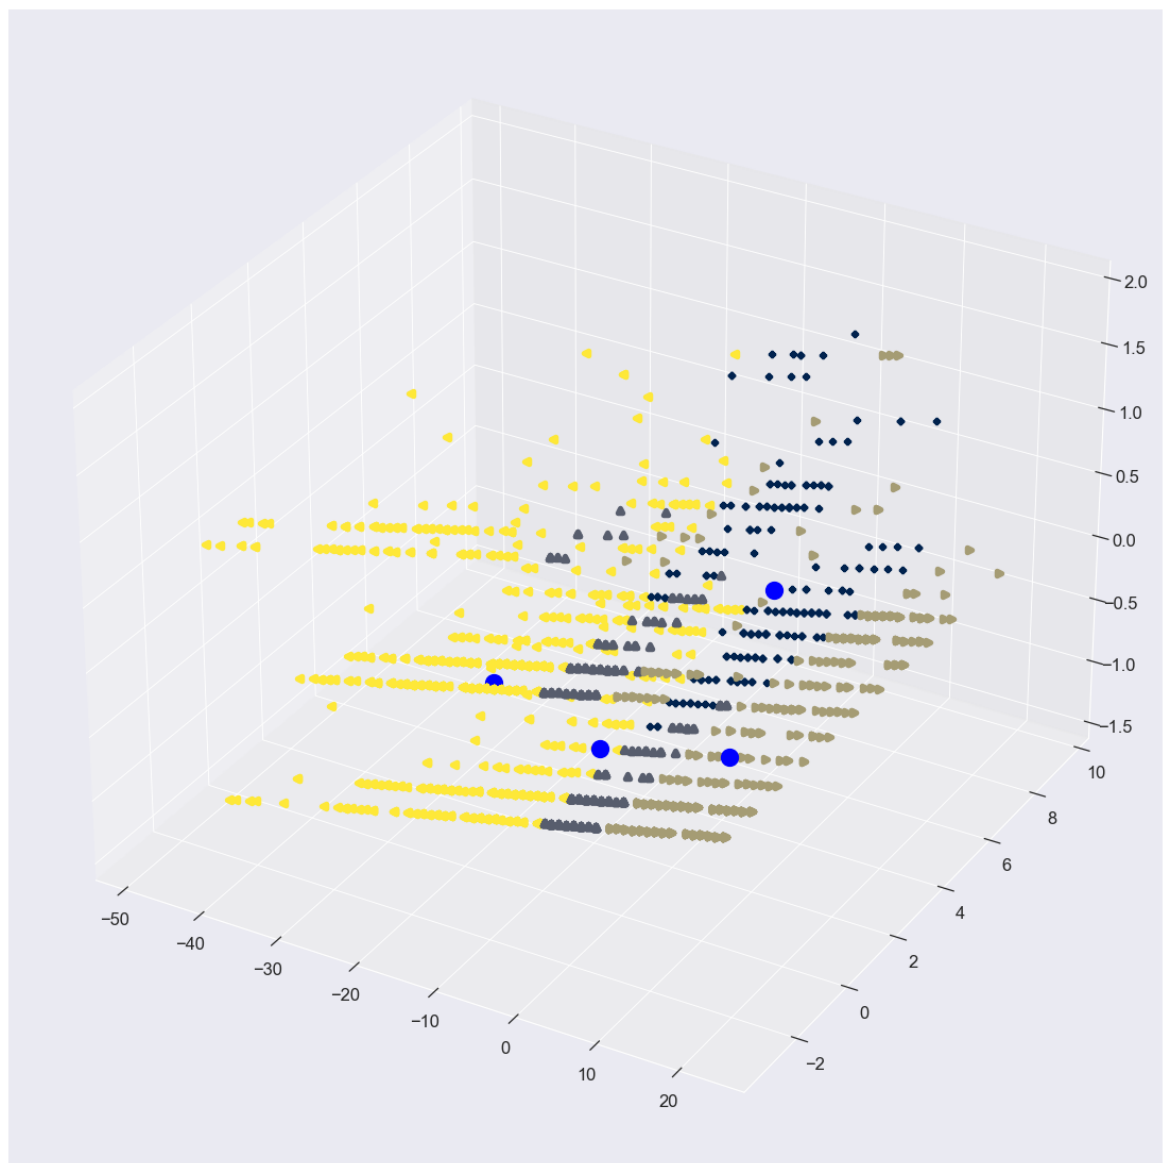
In [1106]: ▶ # Apply model fit on the training data
train_x = train.drop('Segmentation', axis=1)
# Reduce the number of deciding columns
train_x = train[['Age', 'Work_Experience', 'Spending_Score']].values
train_y = train.Segmentation
train_y = pd.Categorical(train_y, categories=['A', 'B', 'C', 'D'], ordered=True).c

# Apply PCA
updated_x = principalComponentAnalysis(train_x, 3)

# Reduce the number of tolerance
train_kmeans = k_means(k=4, tol=0.01)
train_kmeans.model_fit(updated_x)

# Visualize identified k-means clusters
plot_graph(train_kmeans.centroids, train_kmeans.classifications, k=4)

```



```
In [1107]: ▶ # Predict k-means Labels
out_y = train_kmeans.model_predict(updated_x)

# Calculate the silhouette score
score = silhouette_score(updated_x, out_y, metric='euclidean')

print(f'Silhouette Score: {score}%)
```

Silhouette Score: 0.4204211821081394%

### ***Apply model to test data***




```
In [1108]: # Prepare data for applying model

test_x = test.drop('Segmentation', axis=1)
test_x = test[['Age', 'Graduated', 'Work_Experience', 'Spending_Score']].values
test_y = test.Segmentation
test_y = pd.Categorical(test_y, categories=['A', 'B', 'C', 'D'], ordered=True).codes

# Apply PCA to reduce dimensionality
updated_test_x = principalComponentAnalysis(test_x, 3)

# Apply model fit on the testing data
test_kmeans = k_means(k=4)
test_kmeans.model_fit(updated_test_x)

# Visualize the identified k-means clusters
plot_graph(test_kmeans.centroids, test_kmeans.classifications, k=4)
```



```
In [1109]: ▶ # Predict k-means Labels  
out_y = test_kmeans.model_predict(updated_test_x)  
  
# Generate the silhouette score  
score = silhouette_score(updated_test_x, out_y, metric='euclidean')  
  
print(f'Silhoutte Score: {score}%)
```

Silhoutte Score: 0.4486475161692587%

In [ ]: ▶

# DBScan

Author: Jinrong

In [352]: ▶ data.head()

Out[352]:

	Ever_Married	Age	Graduated	Work_Experience	Spending_Score	Family_Size	Var_1	Segment
0	1	61	1	1.0	2	3.0	5	
1	1	63	1	0.0	2	5.0	5	
0	1	39	1	0.0	1	3.0	5	
0	0	23	0	1.0	0	4.0	5	
0	0	18	0	7.0	0	4.0	5	

In [353]: ▶ df\_clean = data.drop(columns=['ID'])  
df\_clean.head()

Out[353]:

	Gender	Ever_Married	Age	Graduated	Work_Experience	Spending_Score	Family_Size	Var_1
0		1	61	1	1.0	2	3.0	5
1		1	63	1	0.0	2	5.0	5
0		1	39	1	0.0	1	3.0	5
0		0	23	0	1.0	0	4.0	5
0		0	18	0	7.0	0	4.0	5

```
In [354]: df_clean.dtypes
```

```
Out[354]: Gender                int8
Ever_Married                 int8
Age                          int64
Graduated                    int8
Work_Experience              float64
Spending_Score               int8
Family_Size                  float64
Var_1                        int8
Segmentation                 object
Artist                       uint8
Doctor                       uint8
Engineer                     uint8
Entertainment                uint8
Executive                     uint8
Healthcare                   uint8
Homemaker                    uint8
Lawyer                       uint8
Marketing                     uint8
dtype: object
```

```
In [355]: cat_cols = df_clean.select_dtypes('object').columns
num_cols = df_clean.select_dtypes('float64').columns
df_clean[cat_cols] = df_clean[cat_cols].apply(lambda x: x.fillna(x.value_counts().index[0]))
df_clean[num_cols] = df_clean[num_cols].apply(lambda x: x.fillna(x.mean()))
```

```
In [356]: df_clean.isna().mean()
```

```
Out[356]: Gender                0.0
Ever_Married                 0.0
Age                          0.0
Graduated                    0.0
Work_Experience              0.0
Spending_Score               0.0
Family_Size                  0.0
Var_1                        0.0
Segmentation                 0.0
Artist                       0.0
Doctor                       0.0
Engineer                     0.0
Entertainment                0.0
Executive                     0.0
Healthcare                   0.0
Homemaker                    0.0
Lawyer                       0.0
Marketing                     0.0
dtype: float64
```

```
In [357]: df_dummy = pd.get_dummies(df_clean.drop(columns='Segmentation'), drop_first=True)
df_dummy.head()
```

Out[357]:

	Gender	Ever_Married	Age	Graduated	Work_Experience	Spending_Score	Family_Size	Value
0	0	1	61	1	1.0	2	3.0	
1	1	1	63	1	0.0	2	5.0	
2	0	1	39	1	0.0	1	3.0	
3	0	0	23	0	1.0	0	4.0	
4	0	0	18	0	7.0	0	4.0	



```

In [358]: class DBSCAN():
    def __init__(self, eps, min_samples=5):
        self.eps = eps # maximum distance between neighbor
        self.min_samples = min_samples # minimum number of neighbors

    # this function calculate distance between all pairs
    def find_distances(self, X):
        return ((X[:,None,:] - X[None,:,:])**2).sum(axis=2)**0.5

    # return the indices of neighbors for specific observation
    def find_neightbors(self, i):
        return set(np.where(self.D[i] < self.eps)[0])

    # find clusters
    def fit(self, X):

        # calculate all pairwise distances
        self.D = self.find_distances(X)

        m = X.shape[0]
        labels_ = [-1]*m
        c = -1

        # go over each observation
        for i in range(m):

            # if it is undefined
            if labels_[i] == -1:

                # find all neighbors of the current node
                neighbors = self.find_neightbors(i)

                # if there are few than min_samples neighbors
                # Label the node as noise
                if len(neighbors) < self.min_samples:
                    labels_[i] = -2
                    continue

                c += 1

                # assgin it to a new cluster
                labels_[i] = c

                # Label all neighbors to the same cluster
                for x in neighbors:
                    labels_[x] = c

                # go over each neighbor
                while neighbors:
                    q = neighbors.pop()
                    labels_[q] = c

                    # find neighbor of neighbors
                    new = self.find_neightbors(q)
                    # if new neighbor are within the limit
                    if len(new) > self.min_samples:
                        # add new neighbors to the list

```

```

        for x in new:
            if x not in neighbors and labels_[x] in {-1, -2}:
                neighbors.add(x)
    # store the final results
    self.labels_ = labels_

```

```

In [359]: ▶ sc = StandardScaler()

X_scaled = df_dummy.copy()
X_scaled[:] = sc.fit_transform(df_dummy)

```

```

In [360]: ▶ eps_val = np.arange(4, 13, 1)
n_cluster = []
for eps in eps_val:
    print(eps)
    dbscan = DBScan(eps=eps)
    dbscan.fit(X_scaled.values)
    n_cluster.append(len(set(dbscan.labels_)))

```

```

4
5
6
7
8
9
10
11
12

```

```

In [361]: ▶ print(eps_val)
print(n_cluster)


```

```

[ 4  5  6  7  8  9 10 11 12]
[8, 3, 2, 1, 1, 1, 1, 1, 1]

```

```
In [362]: plt.plot(eps_val, n_cluster, 'bo-')
plt.xlabel('Epsilon')
plt.ylabel('Number of Clusters');
```



```
In [363]: dbscan = DBScan(eps=4.35)
dbscan.fit(X_scaled.values)
```

```
In [364]: pd.Series(dbscan.labels_).value_counts()
```

```
Out[364]: 1    7613
0     652
2     325
3     229
dtype: int64
```

```
In [365]: pd.crosstab(np.array(dbscan.labels_), data['Segmentation'])
```

```
Out[365]:
```

	Segmentation	A	B	C	D
row_0	0	245	149	164	94
1	3285	1355	1508	1465	
2	138	24	29	134	
3	102	44	19	64	

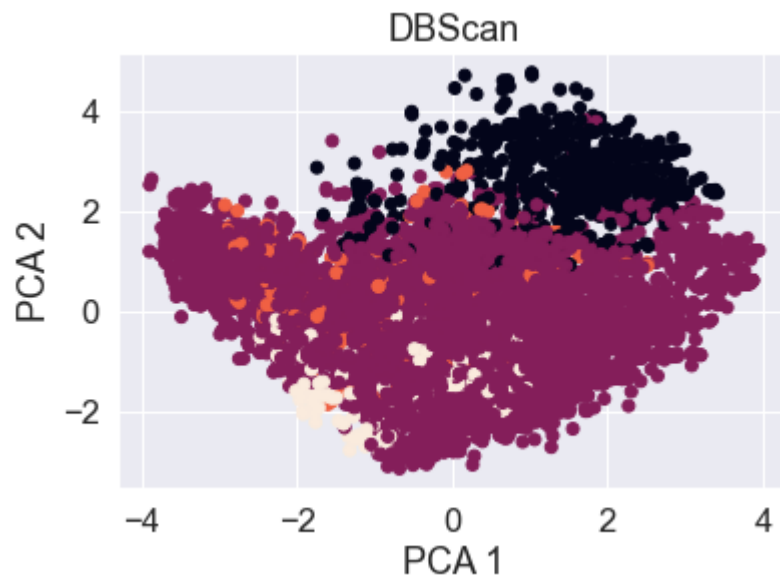
```
In [366]: adjusted_rand_score(data['Segmentation'], dbscan.labels_)
```

```
Out[366]: 0.004704905550590988
```

```
In [367]: X_pca = pca.fit_transform(X_scaled)
```



```
In [368]: ▶ plt.scatter(X_pca[:,0], X_pca[:,1], c=dbscan.labels_)  
plt.title('DBScan')  
plt.xlabel('PCA 1')  
plt.ylabel('PCA 2');
```



```
In [ ]: ▶
```

```
In [ ]: ▶
```

## KNN Algorithm

```
In [118]: ▶ # apply model fit on the training data
X = correlation_data.drop('Segmentation', axis=1)
X = correlation_data[['Age', 'Graduated', 'Work_Ex', 'Spending_Score']]
X
```

Out[118]:

	Age	Graduated	Work_Ex	Spending_Score
0	61	1	1.0	2
1	63	1	0.0	2
2	39	1	0.0	1
3	23	0	1.0	0
4	18	0	7.0	0
...	...	...	...	...
10690	43	1	0.0	1
10691	31	1	1.0	0
10692	22	0	1.0	0
10693	66	1	0.0	1
10694	43	1	1.0	0

8819 rows × 4 columns

```
In [105]: ▶ # #cols = ['Segmentation', 'Artist', 'Doctor', 'Engineer', 'Entertainment', '
# X = correlation_data.drop(['Segmentation', 'Artist', 'Doctor', 'Engineer',
# X
```

```
In [120]: y = correlation_data['Segmentation']  
y
```

```
Out[120]: 0      2  
          1      2  
          2      2  
          3      3  
          4      3  
          ..  
        10690    2  
        10691    3  
        10692    3  
        10693    0  
        10694    1  
          Name: Segmentation, Length: 8819, dtype: int8
```

```
In [107]: # plt.scatter(X[:,0], X[:,1], marker="o", c=y, s=100, cmap="plasma")  
# import matplotlib.pyplot as plt  
# ax.scatter(correlation_data[2], correlation_data[3], correlation_data[4], m  
  
# plt.show()
```

```
In [121]: # split X and y into training and testing sets  
  
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, ra
```

In [122]: `X_train, X_test, y_train, y_test`

```
Out[122]: (
  Age  Graduated  Work_Ex  Spending_Score
10031   50         0      1.0             1
607     46         1      7.0             0
7031    35         0      1.0             0
4774    37         1      8.0             1
10203   58         1      0.0             1
...     ...         ...      ...           ...
5440    67         1      9.0             0
9546    49         1      1.0             2
6000    32         0      0.0             0
4017    79         1      1.0             2
3397    42         1      0.0             1
```

```
[7055 rows x 4 columns],
  Age  Graduated  Work_Ex  Spending_Score
6433   65         1      0.0             1
455    53         1      2.0             0
9929   73         0      0.0             2
4105   29         0      5.0             0
9748   43         1      1.0             0
...     ...         ...      ...           ...
27     47         1      0.0             2
8101   52         1      1.0             0
8615   56         1      0.0             2
681    35         1      9.0             0
10642  42         1      1.0             1
```

```
[1764 rows x 4 columns],
10031    0
607      0
7031     0
4774     2
10203    2
..
5440     1
9546     2
6000     1
4017     1
3397     2
Name: Segmentation, Length: 7055, dtype: int8,
6433     1
455      0
9929     2
4105     0
9748     0
..
27       0
8101     2
8615     1
681      3
10642    2
Name: Segmentation, Length: 1764, dtype: int8)
```

In [123]: `# check the shape of X_train and X_test`

```
X_train.shape, X_test.shape
```

Out[123]: ((7055, 4), (1764, 4))

In [124]: `# Changing the index of the records into range`

```
X_train.index=range(len(X_train))
y_train.index=range(len(X_train))
X_test.index=range(len(X_test))
y_test.index=range(len(y_test))
```

## K nearest neighbours by Sorting the Euclidean distance

In [125]: `def nearest_Neighbours(X_train,y_train,X_test,K):`

```
    dist=[]
    for i in range(len(X_train)):
        #initialize distance = 0
        euclidean_Dist=0
        for j in range(len(X_train.columns)):
            #sum of total distance
            euclidean_Dist+=round(np.sqrt(pow((X_train.iloc[i,j]-X_test[j
            dist.append((euclidean_Dist,i,y_train.iloc[i]))
            #sort in ascending order of distance & select top k nearest distances
            dist = sorted(dist, key=lambda z: z[0])[0:K]
    return dist
```

## Predicting the new data point

```
In [126]: #Predicting the label of the new piece of data based in k-nearest neighbours

def knn_prediction(X_train,y_train,X_test,K):
    neighbours=[]
    pred_outcome=[]
    for i in range(len(X_test)):
        neighbours.append(nearest_Neighbours(X_train,y_train,X_test.iloc[i,:])
    for i in neighbours:
        top_neighbours = {}
        for j in i:
            #list of distances of top k-neighbours
            if j[-1] in top_neighbours.keys():
                top_neighbours[j[-1]]=top_neighbours[j[-1]]+1
            else:
                top_neighbours[j[-1]]=1
        pred_outcome.append(sorted(top_neighbours,key=top_neighbours.get,reverse=True))
    return pred_outcome #return the label
```

## Accuracy calculation of predicted data point

```
In [127]: # Accuracy of correctly predicted
def knn_getAccuracy(actual,predicted):
    correctly_pred=0
    for i in range(len(predicted)):
        if predicted[i]==actual[i]:
            correctly_pred+=1
    return round((correctly_pred/len(actual))*100,2)
```

## Accuracy of Model

```
In [128]: # Accuracy of predicted species
output=knn_prediction(X_train,y_train,X_test,100)

knn_getAccuracy(y_test,output)
```

Out[128]: 48.36

## Checking using KNeighborsClassifier

```
In [129]: # Packages  
%matplotlib notebook  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
from sklearn.model_selection import train_test_split  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.metrics import accuracy_score
```

```
In [130]: # split X and y into training and testing sets  
  
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, ra
```

```
In [131]: knn = KNeighborsClassifier(n_neighbors = 5) #setting up the KNN model to use  
knn.fit(X_train, y_train) #fitting the KNN
```

Out[131]: KNeighborsClassifier()

```
In [150]: #Checking performance on the training set  
print('Accuracy of K-NN classifier on training set: {:.2f}'.format(knn.score(  
#Checking performance on the test set  
print('Accuracy of K-NN classifier on test set: {:.2f}'.format(knn.score(X_te
```

Accuracy of K-NN classifier on training set: 0.53  
Accuracy of K-NN classifier on test set: 0.44

```
In [182]: print("Preliminary model score:")  
print(knn.score(X_test,y_test))
```

Preliminary model score:  
0.4387755102040816

```
In [183]: no_neighbors = np.arange(1, 9)  
train_accuracy = np.empty(len(no_neighbors))  
test_accuracy = np.empty(len(no_neighbors))
```

```
In [184]: for i, k in enumerate(no_neighbors):  
    # We instantiate the classifier  
    knn = KNeighborsClassifier(n_neighbors=k)  
    # Fit the classifier to the training data  
    knn.fit(X_train,y_train)  
  
    # Compute accuracy on the training set  
    train_accuracy[i] = knn.score(X_train, y_train)  
  
    # Compute accuracy on the testing set  
    test_accuracy[i] = knn.score(X_test, y_test)  
    # Visualization of k values vs accuracy  
  
plt.title('k-NN: Varying Number of Neighbors')  
plt.plot(no_neighbors, test_accuracy, label = 'Testing Accuracy')  
plt.plot(no_neighbors, train_accuracy, label = 'Training Accuracy')  
plt.legend()  
plt.xlabel('Number of Neighbors')  
plt.ylabel('Accuracy')  
plt.show()
```

<IPython.core.display.Javascript object>





